

R15

Importing data

Covered in R15

- Overview
- Using datasets in packages
- Importing data frames from other programs
- Importing data from CSV files
- Importing data from Excel files and elsewhere

1 Overview

In the last chapter we looked at the data frame, which is going to be the standard object used for most data analysis.

In this chapter we look at importing data from various sources into R.

There are three fundamental ways of getting data into R:

- manually
- using datasets included in R's packages
- importing it from another program such as Excel.

We have already covered how to manually enter data into vectors or data frames in previous chapters. So this chapter will cover the other two methods.

2 Using datasets in packages

Inbuilt datasets

It may be helpful to review Chapter 8 before reading this section.

We're going to look at the datasets package which, unsurprisingly, contains a variety of datasets.

To find out more about the contents of this package we can click on its name in the Packages window.

Name	Description	Version
boot	Bootstrap Functions (Originally by Angelo Canty for S)	1.3-20
class	Functions for Classification	7.3-14
cli	Helpers for Developing Command Line Interfaces	1.0.0
cluster	"Finding Groups in Data": Cluster Analysis Extended Rousseeuw et al.	2.0.7-1
codetools	Code Analysis Tools for R	0.2-15
colorspace	Color Space Manipulation	1.3-2
compiler	The R Compiler Package	3.5.1
crayon	Colored Terminal Output	1.3.4
crosstalk	Inter-Widget Interactivity for HTML Widgets	1.0.0
datasets	The R Datasets Package	3.5.1
digest	Create Compact Hash Digests of R Objects	0.6.15
dplyr	A Grammar of Data Manipulation	0.7.6
fansi	ANSI Control Sequence Aware String Functions	0.3.0

The R Datasets Package 

Documentation for package 'datasets' version 3.5.1

- [DESCRIPTION file](#).

Help Pages

[A](#) [B](#) [C](#) [D](#) [E](#) [F](#) [H](#) [I](#) [J](#) [L](#) [M](#) [N](#) [O](#) [P](#) [Q](#) [R](#) [S](#) [T](#) [U](#) [V](#) [W](#)

datasets-package	The R Datasets Package
-- A --	
ability.cov	Ability and Intelligence Tests
airmiles	Passenger Miles on Commercial US Airlines, 1937-1960

Clicking on the links will take us to the help page on that specific dataset. For example, clicking on rivers gives:

The screenshot shows the R Documentation interface. At the top, there's a menu bar with 'Files', 'Plots', 'Packages', 'Help', and 'Viewer'. Below the menu is a toolbar with icons for back, forward, search, and help. A search bar says 'Find in Topic'. The main area has a title 'R: Lengths of Major North American Rivers'. Underneath, it says 'rivers {datasets}'. On the right side, there's a vertical scroll bar and a section labeled 'R Documentation'.

Description

This data set gives the lengths (in miles) of 141 "major" rivers in North America, as compiled by the US Geological Survey.

Usage

`rivers`

Format

A vector containing 141 observations.

Source

World Almanac and Book of Facts, 1975, page 406.

References

McNeil, D. R. (1977) *Interactive Data Analysis*. New York: Wiley.

Accessing inbuilt datasets

Since these inbuilt datasets are already loaded into R's workspace we can access them by simply typing their name:

```
> rivers
 [1] 735 320 325 392 524 450 1459 135 465 600 330 336 280 315 870
 [16] 906 202 329 290 1000 600 505 1450 840 1243 890 350 407 286 280
 [31] 525 720 390 250 327 230 265 850 210 630 260 230 360 730 600
 [46] 306 390 420 291 710 340 217 281 352 259 250 470 680 570 350
 [61] 300 560 900 625 332 2348 1171 3710 2315 2533 780 280 410 460 260
 [76] 255 431 350 760 618 338 981 1306 500 696 605 250 411 1054 735
 [91] 233 435 490 310 460 383 375 1270 545 445 1885 380 300 380 377
 [106] 425 276 210 800 420 350 360 538 1100 1205 314 237 610 360 540
 [121] 1038 424 310 300 444 301 268 620 215 652 900 525 246 360 529
 [136] 500 720 270 430 671 1770
> |
```

Here R displays the vector of 141 observations (lengths in miles of 141 major rivers in North America).

Datasets from other packages

Other packages often contain datasets.

Let's load up the MASS package (using the tick box in the Packages window).

If we use `library(help=MASS)` or `help(package="MASS")` we can see that the MASS package contains lots of datasets (87 in fact).

Let's take a look at one of these, Cars93, which contains data from the different types of cars sold in the USA in 1993. From the help page we can see it is a data frame with 93 types of cars (rows) and 27 variables (columns):

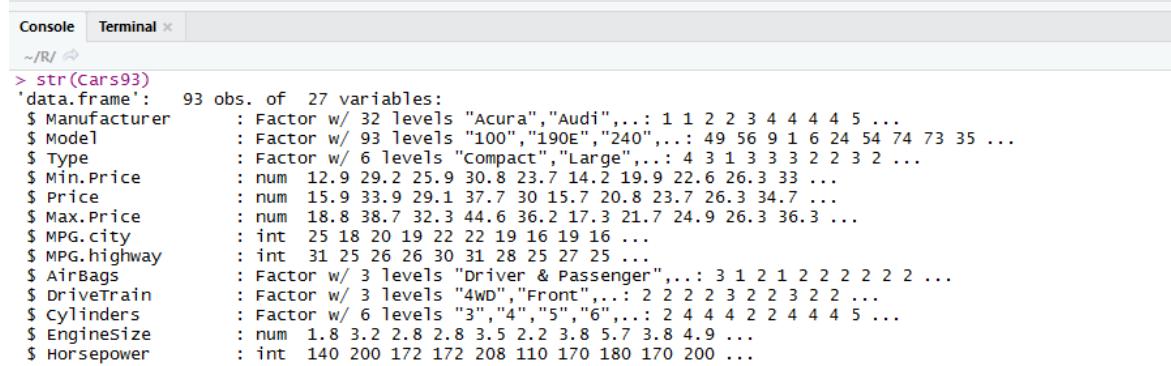
The screenshot shows the R Help browser interface. The top menu bar includes 'Files', 'Plots', 'Packages', 'Help', and 'Viewer'. Below the menu is a toolbar with icons for back, forward, search, and help. The main content area displays the help page for 'Cars93' from the 'MASS' package. The title is 'Data from 93 Cars on Sale in the USA in 1993'. The page is divided into sections: 'Description', 'Usage', 'Cars93', 'Format', and 'Type'. The 'Description' section states: 'The Cars93 data frame has 93 rows and 27 columns.' The 'Usage' section shows the command `Cars93`. The 'Format' section describes the data frame as containing 93 rows and 27 columns. The 'Type' section lists the columns: Manufacturer, Model, and Type. On the right side of the browser, there is a vertical scroll bar and a 'R Documentation' button.

Once we've loaded the package into the workspace we can access the dataset simply by typing its name:

The screenshot shows the R Console interface. The top bar includes 'Console' and 'Terminal'. Below the bar, the command `> Cars93` is entered. The console then displays the first 13 rows of the 'Cars93' data frame. The data is presented as a table with columns: Manufacturer, Model, Type, Min.Price, Price, Max.Price, MPG.city, MPG.highway, AirBags, and DriveTrain. The data shows various car models from different manufacturers like Acura, Audi, BMW, Buick, Cadillac, and Chevrolet, along with their respective details such as engine size, fuel efficiency, and safety features.

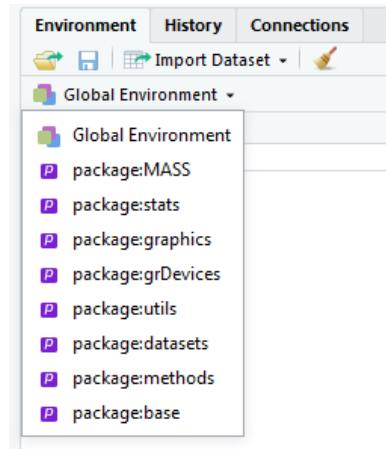
	Manufacturer	Model	Type	Min.Price	Price	Max.Price	MPG.city	MPG.highway	AirBags	DriveTrain
1	Acura	Integra	Small	12.9	15.9	18.8	25	31	None	Front
2	Acura	Legend	Midsize	29.2	33.9	38.7	18	25	Driver & Passenger	Front
3	Audi	90	Compact	25.9	29.1	32.3	20	26	Driver only	Front
4	Audi	100	Midsize	30.8	37.7	44.6	19	26	Driver & Passenger	Front
5	BMW	535i	Midsize	23.7	30.0	36.2	22	30	Driver only	Rear
6	Buick	Century	Midsize	14.2	15.7	17.3	22	31	Driver only	Front
7	Buick	Lesabre	Large	19.9	20.8	21.7	19	28	Driver only	Front
8	Buick	Roadmaster	Large	22.6	23.7	24.9	16	25	Driver only	Rear
9	Buick	Riviera	Midsize	26.3	26.3	26.3	19	27	Driver only	Front
10	Cadillac	DeVille	Large	33.0	34.7	36.3	16	25	Driver only	Front
11	Cadillac	Seville	Midsize	37.5	40.1	42.7	16	25	Driver & Passenger	Front
12	Chevrolet	Cavalier	Compact	8.5	13.4	18.3	25	36	None	Front
13	Chevrolet	Corsica	Compact	11.4	11.4	11.4	25	34	Driver only	Front

Because of the size of this data frame it is not possible to display it easily in the Console. Using the structure command can tell us about the different variables (columns):

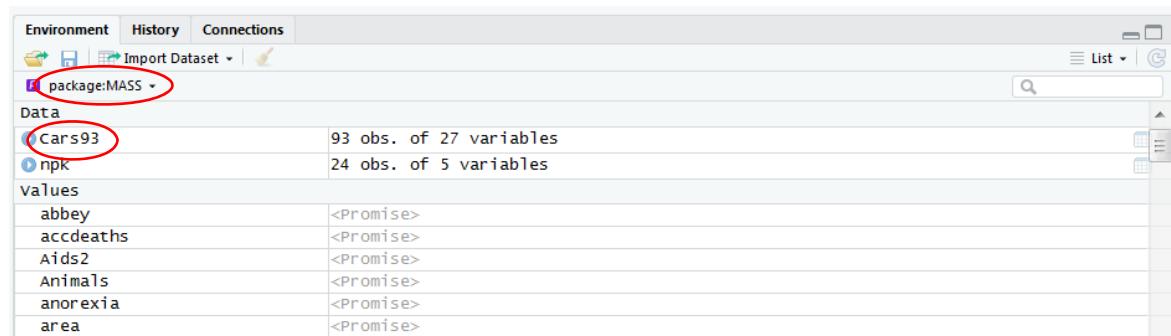


```
Console Terminal ×
~/R/ ↗
> str(Cars93)
'data.frame': 93 obs. of 27 variables:
 $ Manufacturer : Factor w/ 32 levels "Acura","Audi",...: 1 1 2 2 3 4 4 4 4 5 ...
 $ Model        : Factor w/ 93 levels "100","190E",...: 49 56 9 1 6 24 54 74 73 35 ...
 $ Type         : Factor w/ 6 levels "Compact","Large",...: 4 3 1 3 3 3 2 2 3 2 ...
 $ Min.Price    : num 12.9 29.2 25.9 30.8 23.7 14.2 19.9 22.6 26.3 33 ...
 $ Price        : num 15.9 33.9 29.1 37.7 30 15.7 20.8 23.7 26.3 34.7 ...
 $ Max.Price   : num 18.8 38.7 32.3 44.6 36.2 17.3 21.7 24.9 26.3 36.3 ...
 $ MPG.city     : int 25 18 20 19 22 22 19 16 19 16 ...
 $ MPG.highway : int 31 25 26 26 30 31 28 25 27 25 ...
 $ AirBags      : Factor w/ 3 levels "Driver & Passenger",...: 3 1 2 1 2 2 2 2 2 2 ...
 $ DriveTrain   : Factor w/ 3 levels "4WD","Front",...: 2 2 2 2 3 2 2 3 2 2 ...
 $ Cylinders    : Factor w/ 6 levels "3","4","5","6",...: 2 4 4 4 2 2 4 4 4 5 ...
 $ EngineSize   : num 1.8 3.2 2.8 2.8 3.5 2.2 3.8 5.7 3.8 4.9 ...
 $ Horsepower   : int 140 200 172 172 208 110 170 180 170 200 ...
```

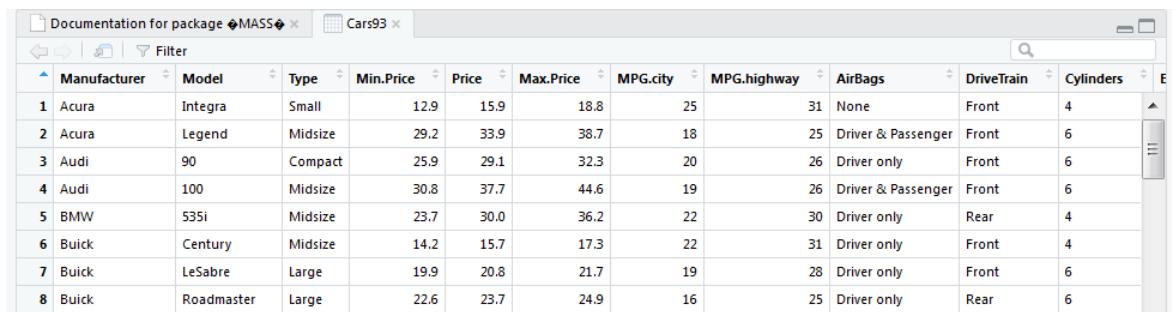
It's easier to view the dataset by selecting it in the Environment window. First select the package by changing *Global Environment* to *package:MASS* using the dropdown menu:



Then select the dataset:



This will open up a new window displaying the dataset in a grid:



The screenshot shows a data grid interface with a header bar containing 'Documentation for package MASS' and 'Cars93'. The grid itself has a header row with columns: Manufacturer, Model, Type, Min.Price, Price, Max.Price, MPG.city, MPG.highway, AirBags, DriveTrain, and Cylinders. Below the header are 8 data rows, each corresponding to a car model from the 'Cars93' dataset. The data includes information such as manufacturer (Acura, Audi, BMW, Buick), model (Integra, Legend, 90, 100, 535i, Century, LeSabre, Roadmaster), type (Small, Midsize, Compact, Midsize, Midsize, Large, Large), and various performance metrics like price and fuel economy.

	Manufacturer	Model	Type	Min.Price	Price	Max.Price	MPG.city	MPG.highway	AirBags	DriveTrain	Cylinders
1	Acura	Integra	Small	12.9	15.9	18.8	25	31	None	Front	4
2	Acura	Legend	Midsize	29.2	33.9	38.7	18	25	Driver & Passenger	Front	6
3	Audi	90	Compact	25.9	29.1	32.3	20	26	Driver only	Front	6
4	Audi	100	Midsize	30.8	37.7	44.6	19	26	Driver & Passenger	Front	6
5	BMW	535i	Midsize	23.7	30.0	36.2	22	30	Driver only	Rear	4
6	Buick	Century	Midsize	14.2	15.7	17.3	22	31	Driver only	Front	4
7	Buick	LeSabre	Large	19.9	20.8	21.7	19	28	Driver only	Front	6
8	Buick	Roadmaster	Large	22.6	23.7	24.9	16	25	Driver only	Rear	6

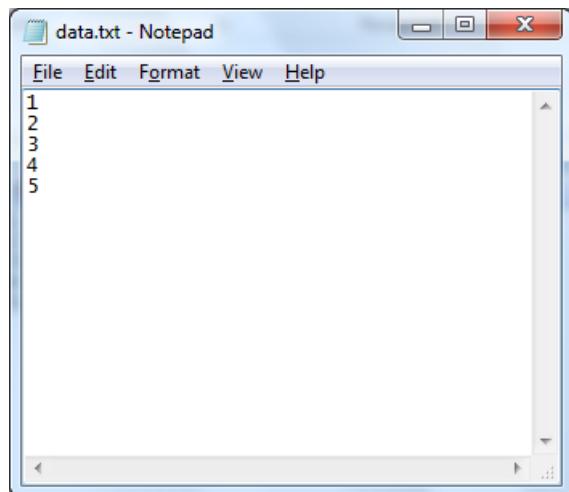
3 Importing data frames from other programs

We can import data from many sources, including:

- text documents
- CSV files
- Excel
- other statistical packages such as SAS or SPSS.

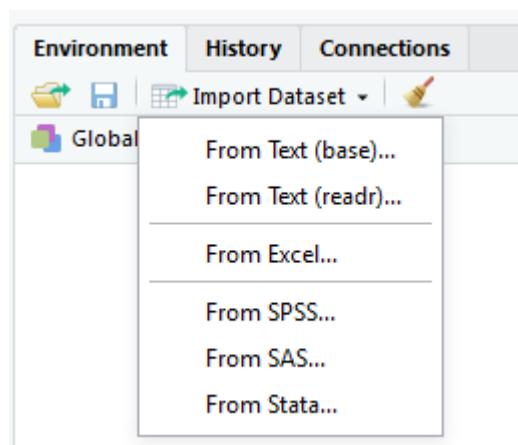
Importing data from text files

Let's create a simple text document. Open up notepad on your machine and enter the numbers 1 to 5 on separate lines. Note that you need to hit return after entering the last number:

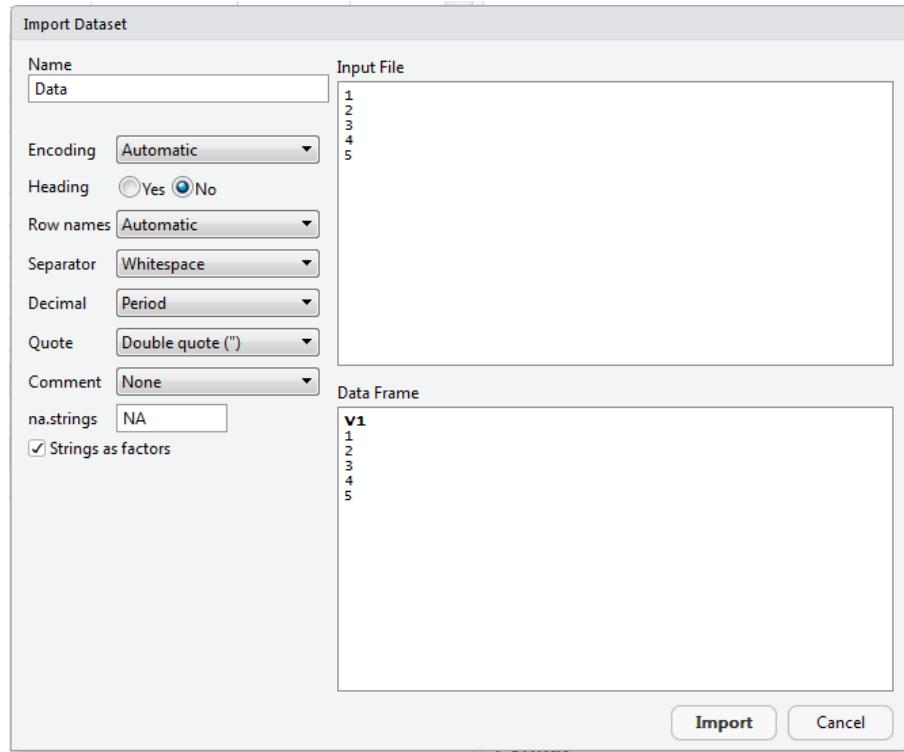


Save this file in your working directory as "Data.txt".

You can import datasets in RStudio using the Environment window. Click on *Import Dataset* and select the first option: *From Text (base)*.

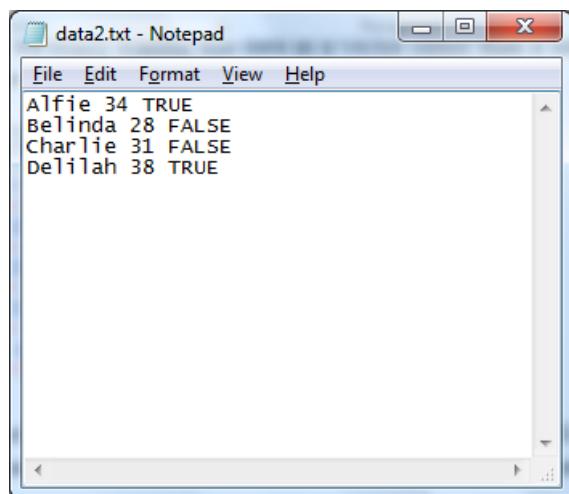


Open the text file you have just created and RStudio will open a window where you can select some options to use for the import (feel free to experiment with the options now or later).

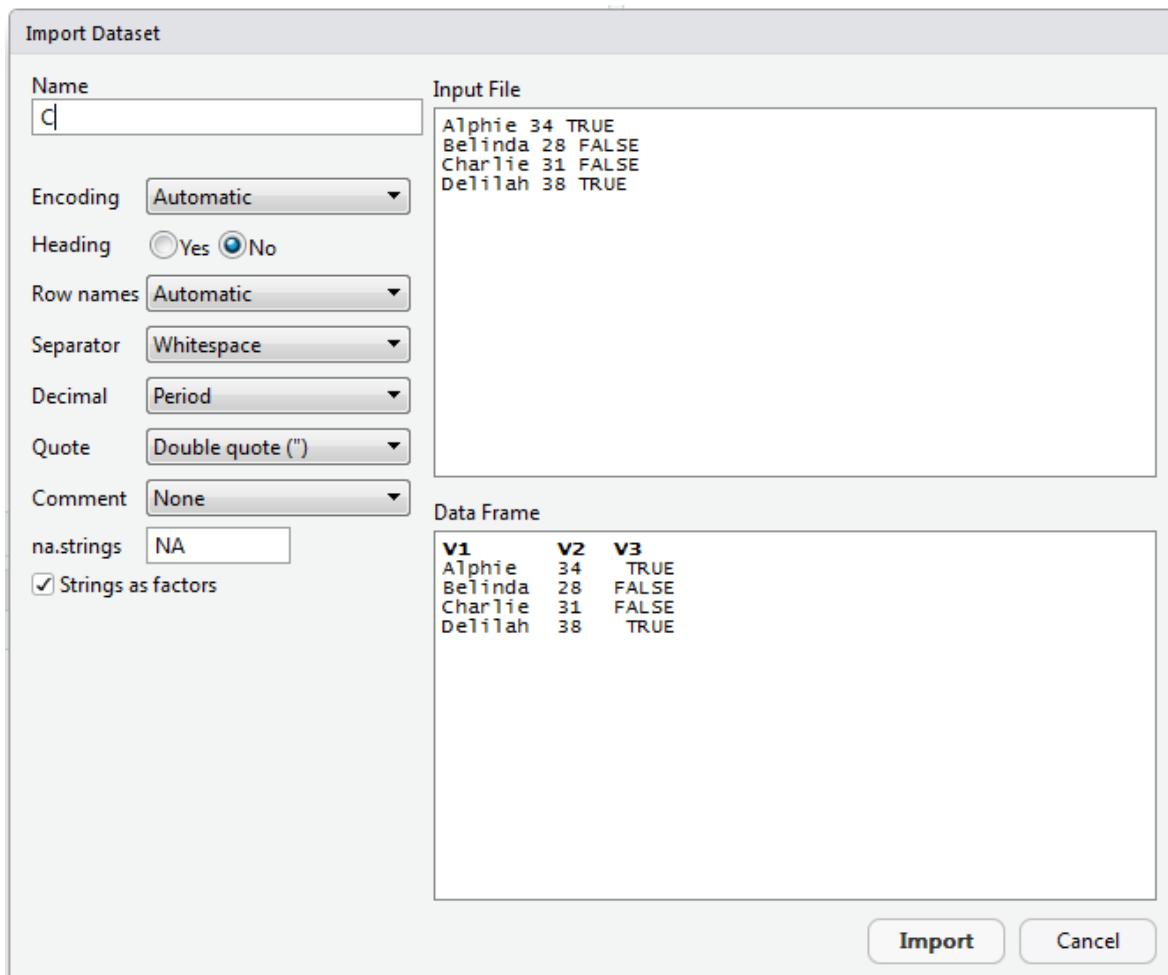


We can see that it has placed a V1 above the numbers. That's because R places the data by default into a data frame and has given the column the name "V1". When you click on *Import* RStudio will display the data in a window and you'll also see the object listed in the Environment window.

Let's now create a three-column dataset in notepad. We can separate the columns with spaces or tabs – both will work. Let's save this dataset as "data2":



We will now load this into object C (by changing the *Name* in the Import window):

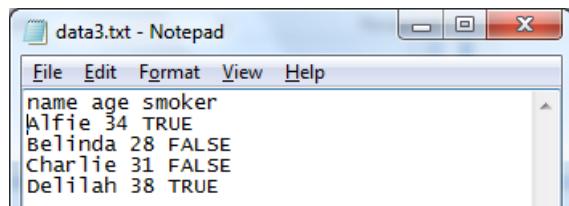


Again, we can see that it is placed in a data frame and by default it has column headings of V followed by the column number.

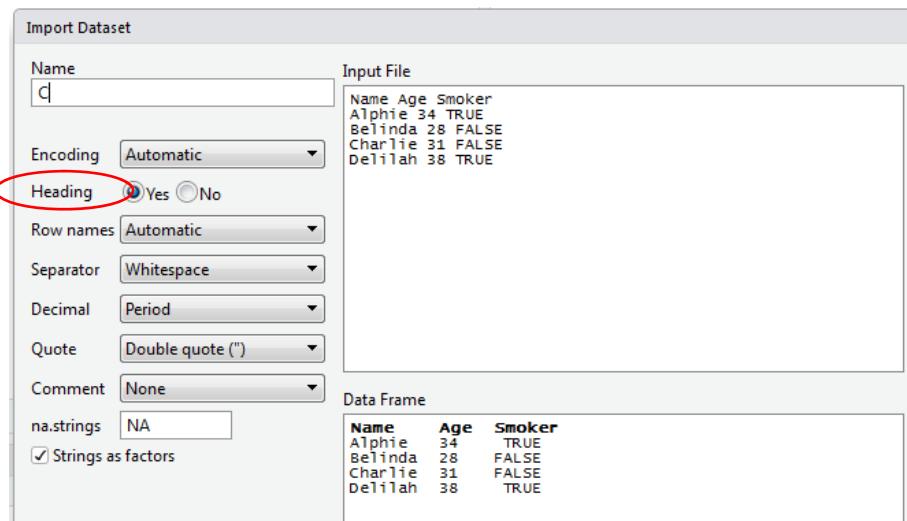
Suppose we want to have column headings "name", "age" and "smoker". One way to do this is to use the colnames (or names) function like we did in Chapter 14. For example:

```
colnames(C) <- c("name", "age", "smoker")
```

An alternative way is to add the column names in the original text document (calling this data3.txt):



When we reach the Import window we can use the *Heading* option.



By default the row names are “1”, “2”, “3”, etc. We could change these using the `rownames` function like we did in Chapter 14. Alternatively, we could use the row names option in the window above and use the first column of data.

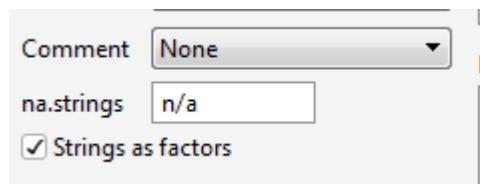
Let's look at the structure of this data frame:

```
> D <- read.table("data3.txt", header=TRUE)
> D
  name age smoker
1 Alfie 34   TRUE
2 Belinda 28 FALSE
3 Charlie 31 FALSE
4 Delilah 38   TRUE
> str(D)
'data.frame': 4 obs. of 3 variables:
 $ name : Factor w/ 4 levels "Alfie","Belinda",...: 1 2 3 4
 $ age  : int 34 28 31 38
 $ smoker: logi TRUE FALSE FALSE TRUE
> |
```

Just like in the previous chapter, R assumes non-numeric data are factors unless we specify otherwise. We can do this by unchecking the *Strings as Factors* box in the Import window.

Missing data

Missing data is often an issue. R uses the logical value NA to tell its functions that the data is missing. However, our text file might use something else. In which case we need to tell R what it is. We do this via the na.strings option in the Import window. For example, we have specified "n/a" and R will replace any occurrences of this with NA:



If you keep an eye on the Console when you import data you will see that R is using the read.table (or possibly the read-delim) function. For example, the last instruction might have read:

```
Data5 <- read.table("~/R/Data3.txt", quote="", comment.char="", na.strings="n/a")
```

If the dataset contained more than one way of indicating missing data, we could adjust this line of code to include the alternatives. For example, here we are looking for "n/a" and "-":

```
Data5 <- read.table("~/R/Data3.txt", quote="", comment.char="", na.strings= c("n/a", "-") )
```

4 Importing data from CSV files

CSV stands for Comma Separated Value and is to Excel what NotePad is to Word. Essentially it's a stripped down excel file that removes all the formatting and separates the values in the cells with, unsurprisingly, commas.

We can import CSV files in the same way as other text files. RStudio will usually recognise the file type and automatically change the *Separator* option to “Comma”. For example:

Import Dataset

Name	CSV	Input File
Encoding	Automatic	Name, Age, Smoker Alphie, 34, TRUE Belinda, 28, FALSE Charlie, 31, FALSE Delilah, 38, TRUE
Heading	<input checked="" type="radio"/> Yes <input type="radio"/> No	
Row names	Automatic	
Separator	Comma	
Decimal	Period	
Quote	Double quote ("")	
Comment	None	
na.strings	NA	
<input checked="" type="checkbox"/> Strings as factors		

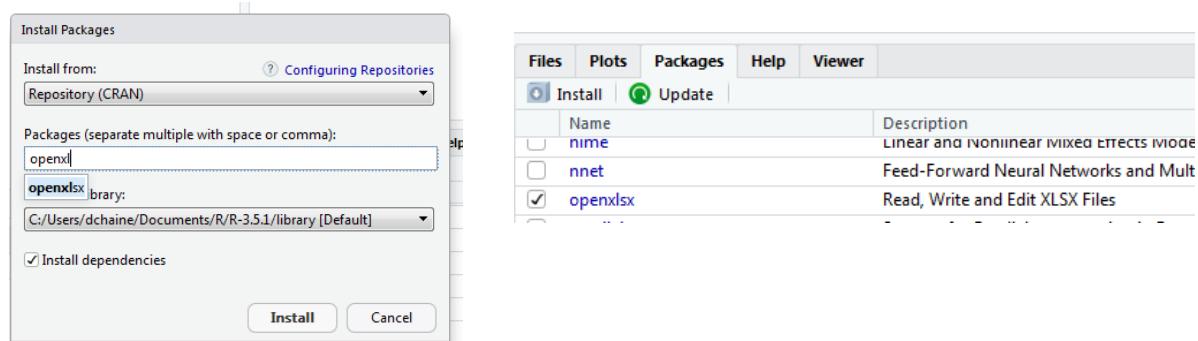
Data Frame

Name	Age	Smoker
Alphie	34	TRUE
Belinda	28	FALSE
Charlie	31	FALSE
Delilah	38	TRUE

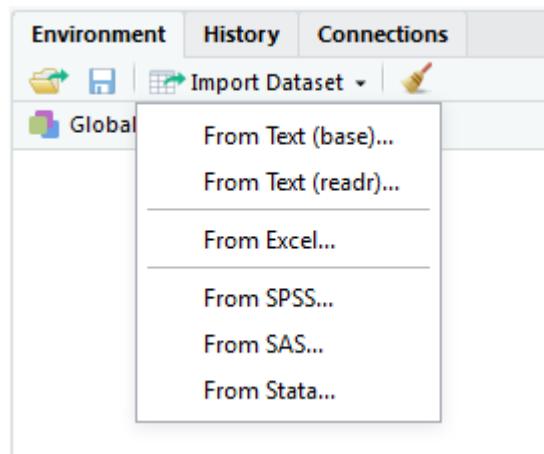
All the other features are the same as before.

5 Importing data from Excel files and elsewhere

To use RStudio's install menu to install Excel files you first need to download and install a package called "openxlsx":



You can select *From Excel* from the import Dataset drop-down menu:



RStudio will then open a window from where you can open the Excel file and select a number of options. Experiment with the options so you can understand what they do.

Importing data from other statistical packages

It is also possible to import data from other statistical packages such as SPSS, SAS or Stata and you can see options for these in the drop-down menu.

You may be prompted to install or update some packages when you first use these options and if you receive an error then search the internet for some help to find out what you might need to install first.

6 Summary

Key terms

Import	Load data stored elsewhere into R
CSV	Comma Separated Value – a format often used to store data sets, with each field of a record separated by a comma

Menus

Import dataset	Located in the Environment window
----------------	-----------------------------------

Key commands

read.table	Imports data into R. View its help for more information on its options.arguments.
------------	-----------------------------------------------------------------------------------

7 Have a go

You will only get proficient at R by practising.

1. Create a dataset in Notepad and import it into R.
2. Create a dataset in Excel, including column and row headings, and import it into R.
3. Alternatively, search the internet for some datasets and load one or two that interest you into R.